

Reinforcement Learning: MDPs, Bellman, and SARSA

BSc Engineering — 3-hour lecture

Carlos Cotrini

2026-05-15

Block 0 — Welcome

Reinforcement Learning in 3 hours

By the end of today you will be able to:

- Read the **SARSA** algorithm
- Explain every line of the pseudocode
- Run SARSA by hand on a 4-state MDP
- Say what SARSA converges to and *why*

Roadmap

Block 1 — Markov Decision Processes · 45 min

 break · 10 min

Block 2 — Returns, value functions, Bellman equations · 50 min

 break · 10 min

Block 3 — ϵ -greedy, SARSA, Q-learning teaser · 50 min

Block 4 — Wrap-up · 10 min


Block 1 — Markov Decision Processes

Hook: a warehouse delivery robot

- 4×4 grid warehouse
- Pick up the **package** at one cell
- Deliver to the **depot** at another
- The floor is **slippery** — the intended move only happens 80 % of the time
- Each step **costs battery**; running out ends the episode

...

The question. What is the best plan?
And how could the robot *learn* one
without being told the rules?



What makes sequential decisions under uncertainty hard?

1. Decisions affect **future states**, not just the immediate outcome.
2. Outcomes are **random**.
3. Reward is **delayed** — the depot is many steps away.

We need a formalism that captures all three. Enter the **Markov Decision Process**.

The Markov property

Intuition. The current state is a *sufficient statistic* of the past — given S_t , knowing $S_{t-1}, S_{t-2}, \dots$ adds nothing.

Formally:

$$\begin{aligned}\Pr[S_{t+1} = s' \mid S_0, A_0, \dots, S_t = s, A_t = a] \\ = \Pr[S_{t+1} = s' \mid S_t = s, A_t = a]\end{aligned}$$

 **Sticky point**

A restriction on the **state**, not on the process. If you need the past, fold it into the state.

Notation we will use today

Random variables, realizations

Symbol	Meaning
S_t, A_t	random state and action at step t
R_{t+1}	random reward after A_t in S_t
s, a, r	realizations / dummy variables

Value functions, hyperparameters


Symbol	Meaning
v_π, q_π	true value functions under π
V, Q	estimates (what algorithms update)
π	policy
$\gamma, \alpha, \varepsilon$	discount, step size, exploration

Sticky point

Reward is indexed $t + 1$, not t . Standard Sutton & Barto convention.

An MDP is a 5-tuple $\langle \mathcal{S}, \mathcal{A}, P, R, \gamma \rangle$

Symbol	Name	Type
$\mathcal{S}$	state space	a set
$\mathcal{A}$	action space	a set
P	transition kernel	$\mathcal{S} \times \mathcal{A} \rightarrow \Delta(\mathcal{S})$
R	reward function	$\mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$
γ	discount factor	scalar in $[0, 1]$



$\mathcal{S}$ — the state space

Everything the agent needs to know to act.

Warehouse robot, partial list:

- Position on grid: $(x, y) \in \{0, 1, 2, 3\}^2$
- Holding the package? $\in \{0, 1\}$
- Battery level: $\in \{0, 1, \dots, B_{\max}\}$

Total: $|\mathcal{S}| = 16 \times 2 \times (B_{\max} + 1)$.

The state is whatever makes the **Markov property** hold for *your* problem.

$\mathcal{A}$ — the action space

What the agent can choose at each step.

Warehouse robot:

$$\mathcal{A} = \{ \text{N, S, E, W, pick-up, drop} \}$$

Action availability can depend on the state — **pick-up** only makes sense on the package cell. We model this either by giving an action-set $\mathcal{A}(s)$ per state, or by allowing the action everywhere and giving illegal cases a large negative reward.

P — the transition kernel


$$P(s' \mid s, a) = \Pr[S_{t+1} = s' \mid S_t = s, A_t = a]$$

A *probability distribution* over next states. Sums to 1 over s' .

Warehouse step (slippery floor):
intended N happens 80 %; left- and right-slips happen 10 % each.

⚠ Warning

In the real world we **do not know** P . That is the whole reason today's lecture exists.



R — the reward function

$R(s, a, s')$ = scalar reward observed when $s \xrightarrow{a} s'$

Warehouse rewards:

- +10 for delivering the package to the depot
- −1 per step (battery cost)
- −100 if the battery dies before delivery




Tip

The reward signal communicates **what** you want — not **how** you want it achieved. *Sutton & Barto*
§3.2

γ — the discount factor

$\gamma \in [0, 1]$ — future rewards are *discounted* by γ per step. Battery **now** > battery in five steps; keeps infinite-horizon sums **finite**.



⚠ Pre-empt

A large γ does **not** mean short-term focus — γ close to 1 makes the agent **patient**.

Trajectories and episodes

A **trajectory** is the sequence the MDP generates:




An **episode** is a trajectory that ends in a *terminal* state.

Warehouse: an episode ends when the package is delivered, or the battery dies.


Policies: deterministic and stochastic

Stochastic: $\pi(a \mid s) = \Pr[A_t = a \mid S_t = s]$ · **Deterministic:** $\pi(s) = a$

Deterministic policy



Stochastic policy (probabilities per action)



We will see soon why we sometimes *want* stochastic policies — to **explore**.

Exercise 1 — identify the MDP components

 Pair work — 3 min

A battery-powered shuttle moves along a one-way 3-stop bus route. At each stop it can **wait** (passengers may board) or **depart**. Each passenger carried gives reward $+5$; each minute waiting costs -1 ; arriving at the depot ends the episode.

1. Write down $\mathcal{S}$ and $\mathcal{A}$.
2. Sketch one valid trajectory of length 4.
3. Write one deterministic policy.

Exercise 1 — solution

$\mathcal{S}$: ($\text{stop} \in \{0, 1, 2\}$, $\text{passengers} \in \mathbb{N}_0$, $\text{minute} \in \mathbb{N}_0$) — the state must contain enough to decide.

$\mathcal{A}$: {wait, depart}.

Trajectory (*one example*):

$(0, 0, 0) \xrightarrow{\text{wait}}_{R=-1} (0, 1, 1) \xrightarrow{\text{depart}}_{R=+5} (1, 0, 2) \xrightarrow{\text{wait}}_{R=-1} (1, 1, 3)$

Deterministic policy: “*depart whenever passenger count ≥ 1 ; otherwise wait.*”

Warning

Common error. Confusing the **trajectory** (what happened) with the **policy** (the decision rule that produced it).



Break (10 min)

Block 2 — Returns, value functions, Bellman

The return G_t

The **undiscounted** return:

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots$$

Two problems:

- May be **infinite** if the episode never ends.
- Treats reward today and reward in 1 000 steps the same.

The **discounted** return:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Why expectations?

The return G_t is random:

- Stochastic environment (P)
- Stochastic policy (π)

Two trajectories from the same starting state can give wildly different G_t .
We need a *single number* per state to compare policies.

Solution. Take the expected value:

$$\text{value} := \mathbb{E}_{\pi}[G_t \mid \text{starting situation}]$$

The subscript π means: “expectation over trajectories generated by following policy π .”

State-value $v_{\pi}(s)$

The expected return if you start in s and follow π forever after:

$$v_{\pi}(s) := \mathbb{E}_{\pi}[G_t \mid S_t = s]$$

Warehouse: “Value of being in this cell, given my current policy.”



Lowercase v_{π} = the **true** function (a property of the MDP and π). Uppercase V = our **estimate** (what algorithms compute).

Action-value $q_\pi(s, a)$

The expected return if you start in s , take action a , and follow π forever after:

$$q_\pi(s, a) := \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a]$$

The first action is **overridden**; everything after that follows π .

⚠ Important

q_π is the **protagonist** of the rest of the lecture. SARSA estimates $Q \approx q_\pi$.

v_π and q_π are connected

The two are linked through the policy:

$$v_\pi(s) = \sum_a \pi(a | s) q_\pi(s, a)$$

In words. The value of a state under π is the policy-weighted average of its action-values at that state.

A consequence: if you have q_π , you have v_π for free.

Bellman expectation: derivation, step 1

Start from the recursive form of the return:

$$G_t = R_{t+1} + \gamma G_{t+1}$$

Plug into the definition of v_π :

$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s]$$

Linearity of expectation:

$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} \mid S_t = s] + \gamma \mathbb{E}_\pi[G_{t+1} \mid S_t = s]$$

Bellman expectation: derivation, step 2

Condition on the action and the next state. By the **Markov property**, $\mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s']$ is just $v_\pi(s')$:

$$\mathbb{E}_\pi[G_{t+1} \mid S_t = s] = \sum_a \pi(a \mid s) \sum_{s'} P(s' \mid s, a) v_\pi(s')$$

Same conditioning for the immediate reward:

$$\mathbb{E}_\pi[R_{t+1} \mid S_t = s] = \sum_a \pi(a \mid s) \sum_{s'} P(s' \mid s, a) R(s, a, s')$$


Combine — and we get the equation on the next slide.

Bellman expectation equation for v_π ★

$$v_\pi(s) = \sum_a \pi(a | s) \sum_{s'} P(s' | s, a) \left[R(s, a, s') + \gamma v_\pi(s') \right]$$

Reading. Pick action under $\pi \rightarrow$ observe r and next state $s' \rightarrow$ value of s is reward plus the discounted value at s' , averaged over everything random.

A **self-consistency** equation: v_π appears on both sides.




Bellman expectation equation for q_π ★★

$$q_\pi(s, a) = \sum_{s'} P(s' | s, a) \left[R(s, a, s') + \gamma \sum_{a'} \pi(a' | s') q_\pi(s', a') \right]$$


⚠ THIS is the equation SARSA samples

The inner structure $r + \gamma q_\pi(s', a')$ for the next (s', a') is **exactly** the (S, A, R, S', A') tuple SARSA collects.



Worked example: compute q_π by hand

3-state chain. Actions **left**, **right**. Deterministic transitions; step cost -1 ; entering terminal T gives $+10$; $\gamma = 0.9$; policy π = always **right**.



Bellman expectation chains back:

$$q_\pi(s_2, \text{right}) = +10$$

$$q_\pi(s_1, \text{right}) = -1 + \gamma q_\pi(s_2, \text{right})$$


$$q_\pi(s_2, \text{left}) = -1 + \gamma q_\pi(s_1, \text{right})$$

(s, a)	q_π
(s_2, right)	$+10$
(s_1, right)	$+8$
(s_2, left)	$+6.2$

Only possible because we knew P, R .

Bellman backups in motion

Bellman backup — sweep 0 / 6



Each sweep propagates the +10 terminal reward one step backwards through the chain — exactly what the Bellman expectation equation prescribes.

Bellman optimality — the dream

The **optimal** value functions:

$$v_*(s) = \max_{\pi} v_{\pi}(s), \quad q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$


There is always a deterministic optimal policy π_* :

$$q_*(s, a) = \sum_{s'} P(s'|s, a) [R + \gamma \max_{a'} q_*(s', a')]$$

Only difference vs. expectation:

$\sum_{a'} \pi(a'|s') \rightarrow \max_{a'}$. The rest of the lecture estimates q_{π} for a gradually-improving policy.

π_* : optimal action per cell




The catch: in practice, P and R are unknown

To solve the Bellman equations directly, we need P and R explicitly. **Real engineering problems** rarely give us either — slip probabilities, contact dynamics, transition kernels are all unknown or approximate.

Three families of solutions:

1. **Estimate P , R from data, then solve** (*Model-based RL — out of scope.*)
2. **Average lots of full episode returns** (*Monte Carlo — high variance, episodic only.*)
3. **Sample one transition and bootstrap** (*Temporal Difference — today.*) ★

Where TD sits — sample × bootstrap



- **DP.** Computes the expectation exactly — needs P and R .
- **MC.** Samples whole episodes — no bootstrapping, but high variance.
- **TD.** Samples one transition, bootstraps off the current estimate.
- **SARSA** lives in the TD corner.

★ Greedy improvement over q_π is model-free

To act greedily, we need $\arg \max_a$.

With $V(s)$:

$$\arg \max_a \sum_{s'} P(s'|s, a) [R + \gamma V(s')]$$

Needs P, R . **Model-based.**

With $Q(s, a)$:

$$\arg \max_a Q(s, a)$$

Just look up the table. **Model-free.**

⚠ Important

This is the reason SARSA estimates q_π instead of v_π .

Exercise 2 — write a Bellman equation

 Pair work — 3 min

The warehouse robot is in cell $(2, 1)$, holding the package, and considers action **N**.

The slip model is the standard 80/10/10. Going **N**:

- Intended (80 %): land in $(2, 2)$
- Slip-left (10 %): land in $(1, 1)$
- Slip-right (10 %): land in $(3, 1)$

Each move costs -1 . The policy π is “*always go N until at the depot.*”

Write the Bellman expectation equation for $q_\pi((2, 1, \text{has-pkg}, \text{batt} = 7), \text{N})$.

Exercise 2 — solution

Apply the boxed Bellman equation. With π deterministic ($\pi(s') = \mathbf{N}$ for every landing state) the inner sum collapses to a single term, and the step cost -1 factors out:

$$q_{\pi}(s, \mathbf{N}) = -1 + \gamma \left[0.8 q_{\pi}(s_N, \mathbf{N}) + 0.1 q_{\pi}(s_W, \mathbf{N}) + 0.1 q_{\pi}(s_E, \mathbf{N}) \right]$$

Watch for


(1) Dropping the $\pi(a'|s')$ weighting (matters once π is stochastic). (2) Double-counting the reward inside the inner sum. (3) Taking $\max_{a'}$ — that's Q-learning, not Bellman expectation for q_{π} .



Break (10 min)

Block 3 — ϵ -greedy and SARSA

Driving home — a TD intuition pump



! Important

MC waits for the trip to end before correcting any estimate. **TD** corrects each estimate the moment the *next* one moves. That is **bootstrapping** — and you don't need an MDP to feel it.

SARSA combines two ideas

To turn the Bellman expectation equation for q_π :

$$q_\pi(s, a) = \sum_{s'} P(s'|s, a) \left[R(s, a, s') + \gamma \sum_{a'} \pi(a'|s') q_\pi(s', a') \right]$$

into something we can **learn from experience**, we make two substitutions.

Idea 1 – Sampling. Replace the expectations $\sum_{s'} P$, $\sum_{a'} \pi$ by single observed samples.

Idea 2 – Bootstrapping. Replace the unknown $q_\pi(s', a')$ by our current estimate $Q(s', a')$.

The next two slides go deep on each idea.

Idea 1 — Sampling

Bellman has $\sum_{s'} P(s'|s, a) \cdot [\dots]$ — an expectation we cannot compute (we don't know P).

But the environment **hands us a sample** s' every time we take action a in state s .

Trick. Replace the expectation by **one** observed transition:

$$\mathbb{E}[X] \approx X \text{ (observed)}$$

Same idea applies to the inner $\sum_{a'} \pi(a'|s')$: just take **one** action a' ourselves.

Idea 2 — Bootstrapping

The Bellman expression contains $q_{\pi}(s', a')$ — a quantity we **don't know**.
But we **have** an estimate $Q(s', a')$. Use it.



Tip

TD updates a guess towards a guess. — *David Silver, UCL Lecture 4*

Yes, the target is a guess. Yes, it's biased. But it gets **better** as Q improves — the algorithm bootstraps itself.

The TD update for Q

Putting both ideas together:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[\underbrace{r + \gamma Q(s', a')}_{\text{TD target}} - \underbrace{Q(s, a)}_{\text{current estimate}} \right]$$

The bracketed quantity is the **TD error**:

$$\delta = \underbrace{r + \gamma Q(s', a')}_{\text{target}} - \underbrace{Q(s, a)}_{\text{estimate}}$$

Reading. Nudge the estimate $Q(s, a)$ a fraction α of the way toward what Bellman would say *if we trusted our current estimates and one observed transition*.

Why we need exploration — the trap

Acting **greedily** with $Q \equiv 0$ initially. Ties broken by action order $N \prec S \prec E \prec W$.

Step 1. All Q s tied at 0 $\rightarrow$ robot picks $N \rightarrow$ step cost $-1 \rightarrow Q(\text{start}, N) = -0.5$.

Step 2. $Q(\text{start}, N) = -0.5$ is now worst. Greedy picks the next-tied action $S \rightarrow$ after the update it too goes negative.

Warning

The robot **never tries** E or **pick-up** long enough to discover reward. Q for those stays at 0, so they “look good” forever — but never get a real chance.


Why we need exploration — the fix

The fix is to occasionally pick a non-greedy action.

That is the role of ε -greedy on the next slide.

ε -greedy

$$A_t = \begin{cases} \arg \max_a Q(S_t, a) & \text{with prob. } 1 - \varepsilon \\ \text{uniform random} & \text{with prob. } \varepsilon \end{cases}$$



ε	behaviour
0	pure exploit — catastrophe (last slide)
1	pure random walk — no learning
0.1	typical default


ϵ schedules

Constant ϵ — simple, but the policy stays *forever* stochastic and never reaches the optimum.

Decaying ϵ ($\epsilon_t = 1/t$, exponential, ...) — explore early, exploit late.

i GLIE

Greedy in the Limit with Infinite Exploration: visit every (s, a) infinitely often AND become greedy in the limit. The formal condition for SARSA to converge to the optimal policy.



SARSA backup diagram



- A chosen ε -greedily from $Q(S, \cdot)$
- R, S' observed from the environment
- A' chosen ε -greedily from $Q(S', \cdot)$

⚠ Important

The chain is the algorithm.

Why “SARSA”?

The five elements of the chain:

$$\left(\textcolor{blue}{S}, \textcolor{orange}{A}, \textcolor{red}{R}, \textcolor{blue}{S}', \textcolor{orange}{A}' \right)$$



Tip

*“This quintuple gives rise to the name **Sarsa** for the algorithm.” – Sutton & Barto, §6.4*

The name IS the mnemonic.

SARSA pseudocode — the whole algorithm

Initialize $Q(s, a)$ arbitrarily for all s, a ; set $Q(\text{terminal}, \cdot) = 0$

Choose hyperparameters $\alpha \in (0, 1]$, $\gamma \in [0, 1]$, $\epsilon \in [0, 1]$

Repeat (for each episode):

$S \leftarrow$ initial state of the episode

$A \leftarrow \epsilon\text{-greedy}(Q, S)$

 Repeat (for each step of episode):

 Take action A ; observe R, S'

$A' \leftarrow \epsilon\text{-greedy}(Q, S')$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'$

$A \leftarrow A'$


 until S is terminal

⚠ Important

Every symbol on this slide has appeared on a previous slide. We have arrived.

Visualising Q on the gridworld

Each cell is **quartered** – one triangle per action $\{N, S, E, W\}$. Number in each triangle = $Q(s, a)$.



Cells near the depot acquire high values first; value **flows backwards** from the depot through the gridworld.

Hand-trace SARSA — setup and first 2 steps

Setup. Robot at $(0, 0)$, no package, full battery. $Q \equiv 0$. $\alpha = 0.5$, $\gamma = 0.9$, $\varepsilon = 0.1$.

Step 1. $A = N$ (greedy, tie-break order). Bumps wall: $S' = (0, 0)$, $R = -1$. $A' = S$ (next tie).

$$Q((0, 0), N) \leftarrow 0 + 0.5 [-1 + 0.9 \cdot 0 - 0] = -0.5$$

Step 2. $A = S$. Lands at $(0, 1)$, $R = -1$. $A' = N$.

$$Q((0, 0), S) \leftarrow 0 + 0.5 [-1 + 0.9 \cdot 0 - 0] = -0.5$$

Hand-trace SARSA — step 3 and the insight

Step 3. Robot at $(0, 1)$. $A = N$ (all Q s here still 0). Lands at $(0, 0)$. $R = -1$. At $(0, 0)$ now, E is the highest action (N, S are at -0.5). $A' = E$.

$$Q((0, 1), N) \leftarrow 0 + 0.5 [-1 + 0.9 \cdot 0 - 0] = -0.5$$

Note

With $Q \equiv 0$ and step cost -1 , **every** step makes the chosen action look worse — so greedy naturally keeps switching, even without ϵ . SARSA is *learning that nothing has worked yet*.

Convergence — the headline result

💡 Sutton & Barto, §6.4

SARSA converges with probability 1 to an optimal policy and action-value function as long as **all state–action pairs are visited an infinite number of times** and the policy **converges in the limit to the greedy policy** (which can be arranged, for example, with ε -greedy policies by setting $\varepsilon = 1/t$).

Two conditions, no proof today:


1. **GLIE.** Visit every (s, a) infinitely often; become greedy in the limit.
2. **Robbins–Monro on α .** $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$.

“*With probability 1*” = almost-sure convergence, in the measure-theoretic sense — same notion you saw in your stochastic processes course.


Showcase — Windy Gridworld

7×10 grid. Wind blows the agent upward by 1 or 2 cells in middle columns, regardless of action. Start **S**, goal **G**. Step cost -1 .

Windy Gridworld — SARSA learned path



Learning curve ($\alpha = 0.5$, $\epsilon = 0.1$)



SARSA with $\alpha = 0.5$, $\epsilon = 0.1$ stabilises at ~ 17 steps; optimal is 15 — the extra 2 are the cost of exploration.

Teaser: Q-learning swaps one symbol

SARSA:

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left[R + \gamma \mathbf{Q(S', A')} - Q(S, A) \right]$$

Target uses the action $\mathbf{A'}$ that the agent will actually take.

Q-learning:

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left[R + \gamma \max_{\mathbf{a'}} \mathbf{Q(S', a')} - Q(S, A) \right]$$

Target uses the **best** possible next action — regardless of what the agent will do.

⚠ Important

One symbol changed. Profoundly different behaviour. Why? Next slide.

SARSA vs Q-learning — the Cliff

Cliff on bottom row: stepping in gives $R = -100$ + reset. Every other step: $R = -1$.



— SARSA — safe — Q-learning — optimal/risky

💡 The punchline

SARSA learns about the policy it's *actually following*. **Q-learning** learns about the policy it *wishes* it could follow.

Exercise 3 — one SARSA update by hand

 Pair work — 3 min

Hyperparameters. $\alpha = 0.1, \gamma = 0.9$.

Q-table (only what you need):

$$Q(s_2, \text{right}) = 5.0 \quad Q(s_3, \text{left}) = 2.0$$

One observed transition:

$$S = s_2, A = \text{right}, R = -1, S' = s_3, A' = \text{left}$$

Compute the new $Q(s_2, \text{right})$.

Exercise 3 — solution

Plug into the SARSA update and simplify:

$$\begin{aligned} Q(s_2, \text{right}) &\leftarrow Q(s_2, \text{right}) + \alpha [R + \gamma Q(s_3, \text{left}) - Q(s_2, \text{right})] \\ &= 5.0 + 0.1 \cdot [-1 + 0.9 \cdot 2.0 - 5.0] \\ &= 5.0 + 0.1 \cdot (-4.2) \end{aligned}$$


$$Q(s_2, \text{right}) = 4.58$$

Common errors

(1) Using $\max_{a'}$ — that's Q-learning. (2) Treating $Q(S, A)$ on the RHS as the just-updated value. (3) Reading $\gamma Q(s', a')$ as $\gamma + Q(s', a')$.

Block 4 — Wrap

From MDP to SARSA in one diagram



Tip

We **defined** the problem (MDP) → **characterised** what we want (q_{π} via Bellman) → **bypassed** the unknown model (sample + bootstrap) → **fixed** exploration (ε -greedy) → **arrived** at SARSA.

What we did *not* cover

Named, not taught — for the curious:

Other tabular methods

- Q-learning (beyond the teaser)
- Expected SARSA, n -step SARSA
- SARSA(λ), eligibility traces

Beyond tabular RL

- Function approximation, Deep RL
- Policy gradients
- Continuous actions, off-policy

Where to read next: Sutton & Barto (free PDF) · David Silver UCL · Stanford CS234.



Tip

You now have the language to read and understand any of these.